

# Practice Lab: CNN for CIFAR10

Training cifar10 Dataset with a CNN model (using tensorflow)

## Outline

- [Packages](#)
- [Loading Data](#)
  - [Dataset overview](#)
- [Modeling](#)
  - [MLP](#)
  - [CNN](#)
- [Evaluating](#)

## 1 - Packages

```
In [53]: import numpy as np
import pandas as pd
import tensorflow as tf
import matplotlib.pyplot as plt
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Conv2D, MaxPool2D, Flatten, BatchNormalization
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras import layers, models
from sklearn.metrics import confusion_matrix, classification_report
from keras.layers import AveragePooling2D
from keras.layers import GlobalMaxPooling2D
from keras.layers import GlobalAveragePooling2D
from tensorflow.keras.utils import to_categorical
import seaborn as sns
from keras.utils import plot_model
```

## 2 - Loading Data

```
In [3]: (x_train, y_train), (x_test, y_test) = tf.keras.datasets.cifar10.load_data()
```

## Dataset overview

The CIFAR-10 dataset consists of 60000 32x32 colour images in 10 classes, with 6000 images per class. There are 50000 training images and 10000 test images.

The dataset is divided into five training batches and one test batch, each with 10000 images. The test batch contains exactly 1000 randomly-selected images from each class. The training batches contain the remaining images in random order, but some training batches may contain more images from one class than another. Between them, the training batches contain exactly 5000 images from each class.

```
In [3]: x_train[0].shape
```

```
Out[3]: (32, 32, 3)
```

```
In [4]: y_train.size
```

```
Out[4]: 50000
```

```
In [5]: x_train.size
```

```
Out[5]: 153600000
```

```
In [6]: np.unique(y_train)
```


```
Out[6]: array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9], dtype=uint8)
```

```
In [37]: labels = ["airplane"  
                 , "automobile"  
                 , "bird"  
                 , "cat"  
                 , "deer"  
                 , "dog"  
                 , "frog"  
                 , "horse"  
                 , "ship"  
                 , "truck"]
```

Plotting

```
In [8]: plt.imshow(x_train[12])  
        print(labels[int(y_train[12])])
```

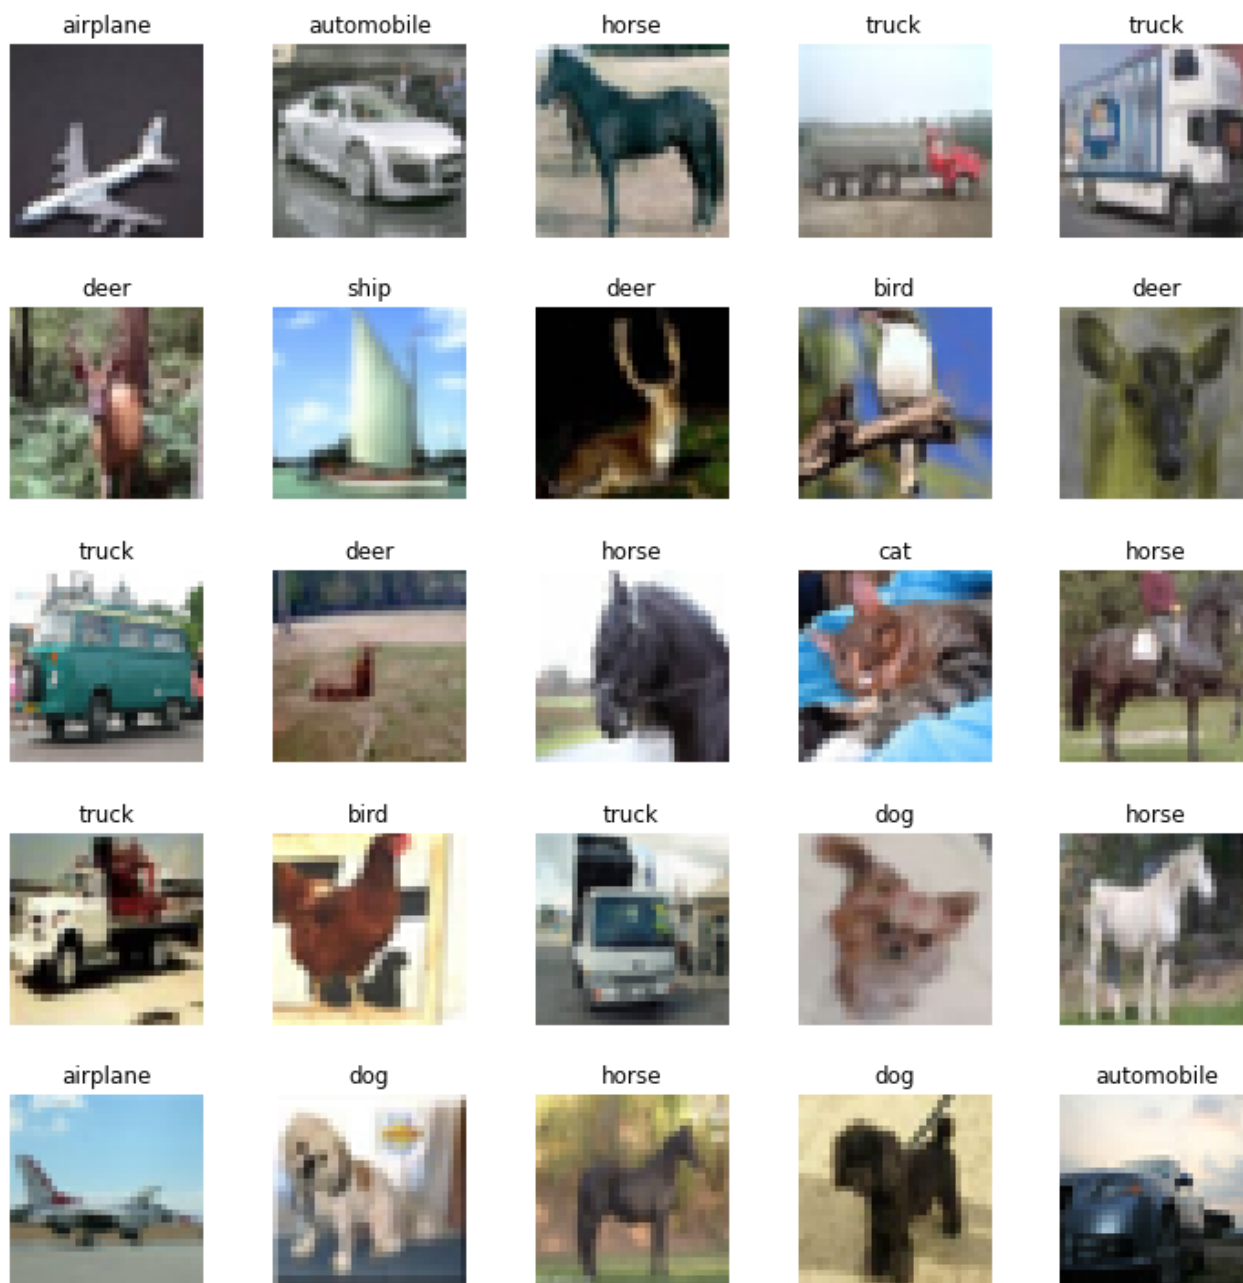
horse



In [38]:

```
import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
X = x_train
fig, axes = plt.subplots(5,5, figsize=(10,10))
fig.tight_layout(pad=0.1)

for i,ax in enumerate(axes.flat):
    random_index = np.random.randint(50000)
    ax.imshow(X[random_index], cmap='gray')
    ax.set_title(labels[int(y_train[random_index])])
    ax.set_axis_off()
```



### 3 - Modeling

First we should Normalize our data

In [4]:

```
x_train = x_train / 255.0
x_test  = x_test  / 255.0
```

In [2]:

```
#Classification report
```

Define classificationReport method :

### Precision

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js

Is the ability of a classifier not to label an instance positive that is actually negative. For each class it is defined as the ratio of true positives to the sum of true and false positives.

## Recall

The ability of a classifier to correctly find all positive instances. For each class, it is defined as the ratio of true positives to the sum of true positives and false negatives.

## F1-score

The F1 score is a weighted harmonic mean of precision and recall such that the best score is 1.0 and the worst is 0.0.

```
In [5]: def ClassificationReport(model):
        y_pred = model.predict(x_test)
        y_pred_classes = [np.argmax(element) for element in y_pred]
        print("Classification Report: \n", classification_report(y_test, y_pred_classes))
```

## 3.1 - simple MLP model :

```
In [15]: mlp = models.Sequential([
        layers.Flatten(input_shape=(32,32,3)),
        layers.Dense(2048, activation='relu'),
        layers.Dense(1024, activation='relu'),
        layers.Dense(10, activation='softmax')
    ])

    mlp.compile(optimizer='SGD',
                loss='sparse_categorical_crossentropy',
                metrics=['accuracy'])

    mlp.fit(x_train, y_train, epochs=5)
```

```
Epoch 1/5
1563/1563 [=====] - 45s 28ms/step - loss: 1.8193 - accuracy: 0.3528
Epoch 2/5
1563/1563 [=====] - 48s 31ms/step - loss: 1.6319 - accuracy: 0.4241
Epoch 3/5
1563/1563 [=====] - 51s 33ms/step - loss: 1.5473 - accuracy: 0.4541
Epoch 4/5
1563/1563 [=====] - 47s 30ms/step - loss: 1.4916 - accuracy: 0.4750
Epoch 5/5
1563/1563 [=====] - 47s 30ms/step - loss: 1.4916 - accuracy: 0.4750
```

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js

1563/1563 [=====] - 48s 31ms/step - loss: 1.4398 - accuracy: 0.4910

Out[15]: <keras.callbacks.History at 0x210f022a9d0>

In [19]: *#Evaluating the MLP model*

In [16]: `mlp.evaluate(x_test,y_test)`

313/313 [=====] - 3s 10ms/step - loss: 1.4760 - accuracy: 0.4643

Out[16]: [1.4760462045669556, 0.4643000066280365]

In [34]: `ClassificationReport(mlp)`

313/313 [=====] - 2s 6ms/step

Classification Report:

	precision	recall	f1-score	support
0	0.58	0.53	0.55	1000
1	0.41	0.82	0.54	1000
2	0.29	0.50	0.37	1000
3	0.44	0.23	0.30	1000
4	0.43	0.31	0.36	1000
5	0.44	0.35	0.39	1000
6	0.51	0.55	0.53	1000
7	0.62	0.45	0.52	1000
8	0.58	0.65	0.61	1000
9	0.66	0.26	0.37	1000
accuracy			0.46	10000
macro avg	0.49	0.46	0.45	10000
weighted avg	0.49	0.46	0.45	10000

Now we use CNN model

## 3.2 - CNN

In [26]:

```

cnn = Sequential([
    layers.Conv2D(filters=32, kernel_size=(3, 3), activation='relu', input_shape=(32, 32, 3)),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(filters=64, kernel_size=(3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(10, activation='softmax')
])

```

In [27]: `cnn.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])`

```
metrics=['accuracy'])
cnn.fit(x_train, y_train, epochs=5)
```

```
Epoch 1/5
1563/1563 [=====] - 24s 15ms/step - loss: 1.4031 - accuracy: 0.4982
Epoch 2/5
1563/1563 [=====] - 27s 17ms/step - loss: 1.0560 - accuracy: 0.6318
Epoch 3/5
1563/1563 [=====] - 27s 17ms/step - loss: 0.9153 - accuracy: 0.6818
Epoch 4/5
1563/1563 [=====] - 26s 17ms/step - loss: 0.8116 - accuracy: 0.7183
Epoch 5/5
1563/1563 [=====] - 26s 17ms/step - loss: 0.7234 - accuracy: 0.7449
```

Out[27]: <keras.callbacks.History at 0x2110a34dfd0>

In [28]: `cnn.evaluate(x_test,y_test)`

Out[28]: 313/313 [=====] - 2s 5ms/step - loss: 0.9776 - accuracy: 0.6693  
[0.9776250123977661, 0.6693000197410583]

In [39]: `cnn.summary()`

Model: "sequential\_3"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_1 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 6, 6, 64)	0
flatten_3 (Flatten)	(None, 2304)	0
dense_9 (Dense)	(None, 128)	295040
dense_10 (Dense)	(None, 10)	1290
=====		
Total params: 315,722		
Trainable params: 315,722		
Non-trainable params: 0		

In [30]: `cnn4 = Sequential([`

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js  
`Layer(conv2d, conv2d_1, max_pooling2d, max_pooling2d_1, flatten_3, dense_9, dense_10, activation='relu', input_shape=(32, 32, 3))`

```

layers.MaxPooling2D((2, 2)),
layers.Conv2D(filters=64, kernel_size=(3, 3), activation='relu'),
layers.MaxPooling2D((2, 2)),
layers.Flatten(),
layers.Dense(128, activation='relu'),
layers.Dense(10, activation='softmax')
])

```

In [38]:

```

cnn64.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
cnn.fit(x_train, y_train, epochs=5)

```

Epoch 1/5

1563/1563 [=====] - 43s 28ms/step - loss: 1.4424 - accuracy: 0.4808

Epoch 2/5

1563/1563 [=====] - 48s 31ms/step - loss: 1.0780 - accuracy: 0.6230

Epoch 3/5

1563/1563 [=====] - 46s 29ms/step - loss: 0.9444 - accuracy: 0.6707

Epoch 4/5

1563/1563 [=====] - 47s 30ms/step - loss: 0.8539 - accuracy: 0.6999

Epoch 5/5

1563/1563 [=====] - 47s 30ms/step - loss: 0.7867 - accuracy: 0.7256

Out[38]: <keras.callbacks.History at 0x2110cffc6d0>

In [40]:

```
cnn64.evaluate(x_test,y_test)
```

313/313 [=====] - 3s 9ms/step - loss: 0.9197 - accuracy: 0.6845  
 Out[40]: [0.9196601510047913, 0.684499979019165]

In [41]:

```
ClassificationReport(cnn64)
```

313/313 [=====] - 3s 9ms/step

Classification Report:

	precision	recall	f1-score	support
0	0.73	0.71	0.72	1000
1	0.82	0.80	0.81	1000
2	0.64	0.50	0.56	1000
3	0.55	0.47	0.51	1000
4	0.65	0.56	0.60	1000
5	0.58	0.59	0.58	1000
6	0.69	0.81	0.74	1000
7	0.71	0.73	0.72	1000
8	0.79	0.82	0.80	1000
9	0.66	0.85	0.74	1000
accuracy				0.68 10000
macro avg				0.68 0.68 0.68 10000

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js



weighted avg	0.68	0.68	0.68	10000
--------------	------	------	------	-------

```
In [48]: cnn32 = Sequential([
    layers.Conv2D(filters=32, kernel_size=(3, 3), activation='relu', input_shape=(32, 32, 3)),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(filters=64, kernel_size=(3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(filters=64, kernel_size=(3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(10, activation='softmax')
])
```


```
In [49]: cnn32.compile(optimizer='adam',
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy'])
cnn32.fit(x_train, y_train, epochs=5)
```

```
Epoch 1/5
1563/1563 [=====] - 26s 16ms/step - loss: 1.5680 - accuracy: 0.4231
Epoch 2/5
1563/1563 [=====] - 31s 20ms/step - loss: 1.2152 - accuracy: 0.5671
Epoch 3/5
1563/1563 [=====] - 31s 20ms/step - loss: 1.0814 - accuracy: 0.6172
Epoch 4/5
1563/1563 [=====] - 29s 19ms/step - loss: 0.9920 - accuracy: 0.6534
Epoch 5/5
1563/1563 [=====] - 29s 18ms/step - loss: 0.9175 - accuracy: 0.6798
Out[49]: <keras.callbacks.History at 0x21108487190>
```

## Types of Pooling Layers

```
In [75]: #Max pooling (32)
```

Max pooling is a pooling operation that selects the maximum element from the region of the feature map covered by the filter. Thus, the output after max-pooling layer would be a feature map containing the most prominent features of the previous feature map.



```
In [76]: cnn_MaxPooling32 = Sequential([
    layers.Conv2D(filters=32, kernel_size=(3, 3), activation='relu', input_shape=(32, 32, 3)),
    layers.MaxPooling2D((2, 2)),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(10, activation='softmax')
])
```

```
In [77]: cnn_MaxPooling32.compile(optimizer='adam',
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy'])
cnn_MaxPooling32.fit(x_train, y_train, epochs=5)
```

Epoch 1/5  
1563/1563 [=====] - 18s 11ms/step - loss: 1.4437 - accuracy: 0.4859  
Epoch 2/5  
1563/1563 [=====] - 20s 13ms/step - loss: 1.1487 - accuracy: 0.5963  
Epoch 3/5  
1563/1563 [=====] - 20s 13ms/step - loss: 1.0188 - accuracy: 0.6433  
Epoch 4/5  
1563/1563 [=====] - 20s 13ms/step - loss: 0.9241 - accuracy: 0.6765  
Epoch 5/5  
1563/1563 [=====] - 20s 13ms/step - loss: 0.8518 - accuracy: 0.7027

Out[77]: <keras.callbacks.History at 0x211371d2df0>

accuracy: 0.7027

```
In [78]: #Max pooling (64)
```

```
In [79]: cnn_MaxPooling64 = Sequential([
    layers.Conv2D(filters=64, kernel_size=(3, 3), activation='relu', input_shape=(32, 32, 3)),
    layers.MaxPooling2D((2, 2)),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(10, activation='softmax')
])
```

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js

```
layers.Dense(10, activation='softmax')
])
```

I use a bigger epoch number for this model because it was too close to average pooling

In [100...

```
cnn_MaxPooling64.compile(optimizer='adam',
                          loss='sparse_categorical_crossentropy',
                          metrics=['accuracy'])
cnn_MaxPooling64.fit(x_train, y_train, epochs=20)
```

```
Epoch 1/20
1563/1563 [=====] - 37s 23ms/step - loss: 0.0662 - accuracy: 0.9780
Epoch 2/20
1563/1563 [=====] - 36s 23ms/step - loss: 0.0594 - accuracy: 0.9801
Epoch 3/20
1563/1563 [=====] - 37s 23ms/step - loss: 0.0585 - accuracy: 0.9817
Epoch 4/20
1563/1563 [=====] - 37s 23ms/step - loss: 0.0549 - accuracy: 0.9824
Epoch 5/20
1563/1563 [=====] - 37s 23ms/step - loss: 0.0588 - accuracy: 0.9812
Epoch 6/20
1563/1563 [=====] - 37s 24ms/step - loss: 0.0580 - accuracy: 0.9817
Epoch 7/20
1563/1563 [=====] - 37s 23ms/step - loss: 0.0614 - accuracy: 0.9810
Epoch 8/20
1563/1563 [=====] - 37s 23ms/step - loss: 0.0563 - accuracy: 0.9824
Epoch 9/20
1563/1563 [=====] - 37s 23ms/step - loss: 0.0490 - accuracy: 0.9842
Epoch 10/20
1563/1563 [=====] - 37s 23ms/step - loss: 0.0460 - accuracy: 0.9849
Epoch 11/20
1563/1563 [=====] - 38s 24ms/step - loss: 0.0570 - accuracy: 0.9823
Epoch 12/20
1563/1563 [=====] - 37s 24ms/step - loss: 0.0515 - accuracy: 0.9836
Epoch 13/20
1563/1563 [=====] - 36s 23ms/step - loss: 0.0426 - accuracy: 0.9872
Epoch 14/20
1563/1563 [=====] - 36s 23ms/step - loss: 0.0555 - accuracy: 0.9834
Epoch 15/20
1563/1563 [=====] - 36s 23ms/step - loss: 0.0439 - accuracy: 0.9867
Epoch 16/20
1563/1563 [=====] - 36s 23ms/step - loss: 0.0482 - accuracy: 0.9867
```

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js

```


Epoch 17/20
1563/1563 [=====] - 36s 23ms/step - loss: 0.0454 - accuracy: 0.9859
Epoch 18/20
1563/1563 [=====] - 37s 24ms/step - loss: 0.0516 - accuracy: 0.9841
Epoch 19/20
1563/1563 [=====] - 36s 23ms/step - loss: 0.0350 - accuracy: 0.9893
Epoch 20/20
1563/1563 [=====] - 36s 23ms/step - loss: 0.0512 - accuracy: 0.9842
Out[100... <keras.callbacks.History at 0x21111bde850>

```

loss: 0.0512 - accuracy: 0.9842

In [72]: `#Average pooling`

Average pooling computes the average of the elements present in the region of feature map covered by the filter. Thus, while max pooling gives the most prominent feature in a particular patch of the feature map, average pooling gives the average of features present in a patch.



```

In [81]: cnn_AveragePooling = Sequential([
    layers.Conv2D(filters=64, kernel_size=(3, 3), activation='relu', input_shape=(32, 32, 3),
    AveragePooling2D(pool_size = 2, strides = 2),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(10, activation='softmax')
])

```

```

In [101... cnn_AveragePooling.compile(optimizer='adam',
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy'])
cnn_AveragePooling.fit(x_train, y_train, epochs=20)

```

```

Epoch 1/20
1563/1563 [=====] - 32s 20ms/step - loss: 0.0583 - accuracy: 0.9804
Epoch 2/20
1563/1563 [=====] - 34s 22ms/step - loss: 0.0597 - accuracy: 0.9809

```

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js

```

Epoch 3/20
1563/1563 [=====] - 34s 22ms/step - loss: 0.0576 - accuracy: 0.9825
Epoch 4/20
1563/1563 [=====] - 34s 22ms/step - loss: 0.0499 - accuracy: 0.9841
Epoch 5/20
1563/1563 [=====] - 34s 22ms/step - loss: 0.0516 - accuracy: 0.9838
Epoch 6/20
1563/1563 [=====] - 34s 22ms/step - loss: 0.0558 - accuracy: 0.9827
Epoch 7/20
1563/1563 [=====] - 34s 22ms/step - loss: 0.0501 - accuracy: 0.9848
Epoch 8/20
1563/1563 [=====] - 36s 23ms/step - loss: 0.0472 - accuracy: 0.9856
Epoch 9/20
1563/1563 [=====] - 34s 22ms/step - loss: 0.0437 - accuracy: 0.9866
Epoch 10/20
1563/1563 [=====] - 34s 22ms/step - loss: 0.0516 - accuracy: 0.9841
Epoch 11/20
1563/1563 [=====] - 34s 22ms/step - loss: 0.0447 - accuracy: 0.9865
Epoch 12/20
1563/1563 [=====] - 35s 22ms/step - loss: 0.0494 - accuracy: 0.9849
Epoch 13/20
1563/1563 [=====] - 38s 24ms/step - loss: 0.0412 - accuracy: 0.9873
Epoch 14/20
1563/1563 [=====] - 39s 25ms/step - loss: 0.0477 - accuracy: 0.9853
Epoch 15/20
1563/1563 [=====] - 36s 23ms/step - loss: 0.0392 - accuracy: 0.9877
Epoch 16/20
1563/1563 [=====] - 34s 22ms/step - loss: 0.0467 - accuracy: 0.9855
Epoch 17/20
1563/1563 [=====] - 39s 25ms/step - loss: 0.0439 - accuracy: 0.9875
Epoch 18/20
1563/1563 [=====] - 39s 25ms/step - loss: 0.0412 - accuracy: 0.9876
Epoch 19/20
1563/1563 [=====] - 39s 25ms/step - loss: 0.0395 - accuracy: 0.9881
Epoch 20/20
1563/1563 [=====] - 33s 21ms/step - loss: 0.0327 - accuracy: 0.9903

```

Out[101... <keras.callbacks.History at 0x21136e45460>

loss: 0.0327 - accuracy: 0.9903

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js

#Global pooling

Global pooling reduces each channel in the feature map to a single value. Thus, an  $n_h \times n_w \times n_c$  feature map is reduced to  $1 \times 1 \times n_c$  feature map. This is equivalent to using a filter of dimensions  $n_h \times n_w$  i.e. the dimensions of the feature map.

```
In [178... cnn_GlobalPooling = Sequential([
    layers.Conv2D(filters=64, kernel_size=(3, 3), activation='relu', input_shape=(32, 32, 3),
    layers.GlobalAveragePooling2D(),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(10, activation='softmax')
])
```

```
In [179... cnn_GlobalPooling.compile(optimizer='adam',
                             loss='sparse_categorical_crossentropy',
                             metrics=['accuracy'])
cnn_GlobalPooling.fit(x_train, y_train, epochs=5)
```

```
Epoch 1/5
1563/1563 [=====] - 15s 10ms/step - loss: 2.0315 - accuracy: 0.2330
Epoch 2/5
1563/1563 [=====] - 17s 11ms/step - loss: 1.7967 - accuracy: 0.3192
Epoch 3/5
1563/1563 [=====] - 20s 13ms/step - loss: 1.7135 - accuracy: 0.3510
Epoch 4/5
1563/1563 [=====] - 21s 13ms/step - loss: 1.6892 - accuracy: 0.3646
Epoch 5/5
1563/1563 [=====] - 18s 12ms/step - loss: 1.6691 - accuracy: 0.3740
Out[179... <keras.callbacks.History at 0x2112101b0a0>
```

accuracy: 0.3740

```
In [102... cnn_AveragePooling.evaluate(x_test,y_test)
```

```
313/313 [=====] - 2s 6ms/step - loss: 4.3350 - accuracy: 0.6109
Out[102... [4.334970951080322, 0.6108999848365784]
```

```
In [103... cnn_MaxPooling64.evaluate(x_test,y_test)
```

```
313/313 [=====] - 2s 6ms/step - loss: 4.2818 - accuracy: 0.6206
Out[103... [4.2817840576171875, 0.6205999851226807]
```

```
In [114... #Let's combine these for the best result
```

```

layers.Conv2D(filters=64, kernel_size=(3, 3), activation='relu', input_shape=(32, 32, 3),
BatchNormalization(),
AveragePooling2D(pool_size = 2, strides = 2),
layers.Conv2D(filters=64, kernel_size=(3, 3), activation='relu', input_shape=(32, 32, 3),
layers.MaxPooling2D((2, 2)),
layers.Flatten(),
layers.Dense(512,activation='relu'),
layers.Dense(10, activation='softmax')
])

```

In [7]:

```

cnn_Final.compile(optimizer='adam',
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])
cnn_Final.fit(x_train, y_train, epochs=15)

```

Epoch 1/15

1563/1563 [=====] - 63s 40ms/step - loss: 1.2560 - accuracy: 0.5557

Epoch 2/15

1563/1563 [=====] - 59s 38ms/step - loss: 0.8956 - accuracy: 0.6872

Epoch 3/15

1563/1563 [=====] - 62s 39ms/step - loss: 0.6932 - accuracy: 0.7587

Epoch 4/15

1563/1563 [=====] - 64s 41ms/step - loss: 0.5263 - accuracy: 0.8137

Epoch 5/15

1563/1563 [=====] - 65s 41ms/step - loss: 0.3675 - accuracy: 0.8719

Epoch 6/15

1563/1563 [=====] - 65s 41ms/step - loss: 0.2514 - accuracy: 0.9128

Epoch 7/15

1563/1563 [=====] - 65s 41ms/step - loss: 0.1922 - accuracy: 0.9323

Epoch 8/15

1563/1563 [=====] - 67s 43ms/step - loss: 0.1448 - accuracy: 0.9492

Epoch 9/15

1563/1563 [=====] - 63s 40ms/step - loss: 0.1282 - accuracy: 0.9571

Epoch 10/15

1563/1563 [=====] - 64s 41ms/step - loss: 0.1182 - accuracy: 0.9610

Epoch 11/15

1563/1563 [=====] - 65s 42ms/step - loss: 0.1042 - accuracy: 0.9659

Epoch 12/15

1563/1563 [=====] - 64s 41ms/step - loss: 0.0903 - accuracy: 0.9704

Epoch 13/15

1563/1563 [=====] - 61s 39ms/step - loss: 0.0918 - accuracy: 0.9708

Epoch 14/15

1563/1563 [=====] - 59s 38ms/step - loss: 0.0771 - accuracy: 0.9740

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js

1563/1563 [=====] - 57s 36ms/step - loss: 0.0868 - accuracy: 0.9724

Out[7]: <keras.callbacks.History at 0x1a418071e20>

In [115... *#First try*

In [111... cnn\_Final.summary()

Model: "sequential\_29"

Layer (type)	Output Shape	Param #
conv2d_47 (Conv2D)	(None, 30, 30, 64)	1792
average_pooling2d_8 (AveragePooling2D)	(None, 15, 15, 64)	0
conv2d_48 (Conv2D)	(None, 13, 13, 64)	36928
max_pooling2d_32 (MaxPooling2D)	(None, 6, 6, 64)	0
flatten_27 (Flatten)	(None, 2304)	0
dense_66 (Dense)	(None, 128)	295040
batch_normalization (Batch Normalization)	(None, 128)	512
dense_67 (Dense)	(None, 10)	1290
Total params: 335,562		
Trainable params: 335,306		
Non-trainable params: 256		

loss: 0.4474 - accuracy: 0.8409

In [116... *#Second try*

In [117... cnn\_Final.summary()

Model: "sequential\_30"

Layer (type)	Output Shape	Param #
conv2d_49 (Conv2D)	(None, 30, 30, 64)	1792
average_pooling2d_9 (AveragePooling2D)	(None, 15, 15, 64)	0
conv2d_50 (Conv2D)	(None, 13, 13, 64)	36928
max_pooling2d_33 (MaxPooling2D)	(None, 6, 6, 64)	0

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js



g2D)

flatten_28 (Flatten)	(None, 2304)	0
dense_68 (Dense)	(None, 256)	590080
dense_69 (Dense)	(None, 10)	2570

```

=====
Total params: 631,370
Trainable params: 631,370
Non-trainable params: 0

```

---

loss: 0.1076 - accuracy: 0.9618

In [127...

#Third Try

In [126...

cnn\_Final.summary()

Model: "sequential\_32"

Layer (type)	Output Shape	Param #
conv2d_57 (Conv2D)	(None, 30, 30, 64)	1792
average_pooling2d_13 (AveragePooling2D)	(None, 15, 15, 64)	0
conv2d_58 (Conv2D)	(None, 13, 13, 64)	36928
max_pooling2d_37 (MaxPooling2D)	(None, 6, 6, 64)	0
flatten_32 (Flatten)	(None, 2304)	0
dense_75 (Dense)	(None, 256)	590080
dense_76 (Dense)	(None, 128)	32896
dense_77 (Dense)	(None, 10)	1290

```

=====
Total params: 662,986
Trainable params: 662,986
Non-trainable params: 0

```

---

loss: 0.1673 - accuracy: 0.9415

In [128...

# 4th try

In [131...

cnn\_Final.summary()

Model: "sequential\_33"

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js

Layer (type)	Output Shape	Param #
conv2d_59 (Conv2D)	(None, 30, 30, 64)	1792
average_pooling2d_14 (AveragePooling2D)	(None, 15, 15, 64)	0
conv2d_60 (Conv2D)	(None, 13, 13, 64)	36928
max_pooling2d_38 (MaxPooling2D)	(None, 6, 6, 64)	0
flatten_33 (Flatten)	(None, 2304)	0
dense_78 (Dense)	(None, 512)	1180160
dense_79 (Dense)	(None, 10)	5130
Total params: 1,224,010		
Trainable params: 1,224,010		
Non-trainable params: 0		

loss: 0.0787 - accuracy: 0.9730

In [136... *# 5th try*

In [137... `cnn_Final.summary()`

Model: "sequential\_35"

Layer (type)	Output Shape	Param #
conv2d_63 (Conv2D)	(None, 30, 30, 64)	1792
average_pooling2d_16 (AveragePooling2D)	(None, 15, 15, 64)	0
conv2d_64 (Conv2D)	(None, 13, 13, 64)	36928
max_pooling2d_40 (MaxPooling2D)	(None, 6, 6, 64)	0
flatten_35 (Flatten)	(None, 2304)	0
dense_83 (Dense)	(None, 1024)	2360320
batch_normalization_1 (Batch Normalization)	(None, 1024)	4096
dense_84 (Dense)	(None, 512)	524800
dense_85 (Dense)	(None, 10)	5130
Total params: 2,321,010		
Trainable params: 2,321,010		
Non-trainable params: 0		

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js

Non-trainable params: 2,048

loss: 0.2206 - accuracy: 0.9224

In [14]: `#Final try`

In [15]: `cnn_Final.summary()`

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
conv2d_2 (Conv2D)	(None, 30, 30, 64)	1792
batch_normalization_1 (Batch Normalization)	(None, 30, 30, 64)	256
average_pooling2d_1 (Average Pooling2D)	(None, 15, 15, 64)	0
conv2d_3 (Conv2D)	(None, 13, 13, 64)	36928
max_pooling2d_1 (Max Pooling2D)	(None, 6, 6, 64)	0
flatten_1 (Flatten)	(None, 2304)	0
dense_2 (Dense)	(None, 512)	1180160
dropout (Dropout)	(None, 512)	0
dense_3 (Dense)	(None, 10)	5130

Total params: 1,224,266

Trainable params: 1,224,138

Non-trainable params: 128

## 4 - Let's evaluate our model

In [9]: `cnn_Final.evaluate(x_test,y_test)`

313/313 [=====] - 5s 15ms/step - loss: 2.2765 - accuracy: 0.6839

Out[9]: [2.2765324115753174, 0.683899998664856]

loss: 2.2765 - accuracy: 0.6839

In [10]: `ClassificationReport(cnn_Final)`

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js

313/313 [=====] - 6s 18ms/step

Classification Report:

	precision	recall	f1-score	support
0	0.74	0.73	0.74	1000
1	0.84	0.80	0.82	1000
2	0.57	0.57	0.57	1000
3	0.46	0.52	0.48	1000
4	0.69	0.57	0.62	1000
5	0.57	0.56	0.57	1000
6	0.70	0.81	0.75	1000
7	0.76	0.70	0.73	1000
8	0.77	0.83	0.80	1000
9	0.78	0.75	0.77	1000
accuracy			0.68	10000
macro avg	0.69	0.68	0.68	10000
weighted avg	0.69	0.68	0.68	10000

As you can see we got a big loss number for testing data

Now we prevent overfitting for our model

```
In [23]: cnn_Final = Sequential([
    layers.Conv2D(filters=64, kernel_size=(3, 3), activation='relu', input_shape=(32, 32, 3)),
    BatchNormalization(),
    AveragePooling2D(pool_size = 2, strides = 2),
    layers.Conv2D(filters=64, kernel_size=(3, 3), activation='relu', input_shape=(32, 32, 3)),
    layers.MaxPooling2D((2, 2)),
    layers.Flatten(),
    layers.Dense(512, activation='relu'),
    Dropout(0.5),
    layers.Dense(10, activation='softmax')
])
```

```
In [24]: cnn_Final.compile(optimizer='adam',
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy'])
```

```
In [25]: history = cnn_Final.fit(x_train, y_train, validation_data=(x_test, y_test), epochs=7, batch_size=128)
```

Epoch 1/7

167/167 [=====] - 48s 284ms/step - loss: 1.6067 - accuracy: 0.4260 - val\_loss: 2.2226 - val\_accuracy: 0.3002

Epoch 2/7

167/167 [=====] - 48s 287ms/step - loss: 1.2128 - accuracy: 0.5703 - val\_loss: 2.1073 - val\_accuracy: 0.4955

Epoch 3/7

167/167 [=====] - 49s 291ms/step - loss: 1.0481 - accuracy: 0.6322 - val\_loss: 1.3993 - val\_accuracy: 0.6226

Epoch 4/7

167/167 [=====] - 49s 291ms/step - loss: 0.9406 - accuracy: 0.6712 - val\_loss: 0.9649 - val\_accuracy: 0.6710

Epoch 5/7

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js [=====] - 49s 290ms/step - loss: 0.8643 - accuracy: 0.6979 - val\_loss: 0.9692 - val\_accuracy: 0.6629

```
Epoch 6/7
167/167 [=====] - 49s 295ms/step - loss: 0.7951 - accuracy: 0.7
228 - val_loss: 0.8960 - val_accuracy: 0.6911
Epoch 7/7
167/167 [=====] - 49s 292ms/step - loss: 0.7357 - accuracy: 0.7
420 - val_loss: 0.9560 - val_accuracy: 0.6789
```

Try with binary cross entropy

```
In [19]: y_categorical_train = to_categorical(y_train, 10)
y_categorical_test = to_categorical(y_test, 10)
```

```
In [26]: cnn_Final.compile(loss="categorical_crossentropy", optimizer="adam", metrics=["accuracy
```

```
In [28]: history = cnn_Final.fit(x_train, y_categorical_train, validation_data=(x_test, y_catego
```

```
Epoch 1/7
167/167 [=====] - 48s 285ms/step - loss: 0.6871 - accuracy: 0.7
561 - val_loss: 0.8966 - val_accuracy: 0.6960
Epoch 2/7
167/167 [=====] - 48s 289ms/step - loss: 0.6258 - accuracy: 0.7
788 - val_loss: 0.8495 - val_accuracy: 0.7120
Epoch 3/7
167/167 [=====] - 48s 289ms/step - loss: 0.5705 - accuracy: 0.7
995 - val_loss: 0.8333 - val_accuracy: 0.7230
Epoch 4/7
167/167 [=====] - 49s 291ms/step - loss: 0.5290 - accuracy: 0.8
132 - val_loss: 0.8315 - val_accuracy: 0.7213
Epoch 5/7
167/167 [=====] - 50s 297ms/step - loss: 0.4872 - accuracy: 0.8
282 - val_loss: 0.8431 - val_accuracy: 0.7282
Epoch 6/7
167/167 [=====] - 50s 297ms/step - loss: 0.4408 - accuracy: 0.8
435 - val_loss: 0.8508 - val_accuracy: 0.7348
Epoch 7/7
167/167 [=====] - 51s 304ms/step - loss: 0.4071 - accuracy: 0.8
548 - val_loss: 0.9347 - val_accuracy: 0.7089
```

```
In [29]: Y_prediction = cnn_Final.predict(x_test)

313/313 [=====] - 4s 11ms/step
```

```
In [30]: Y_prediction.shape
```

```
Out[30]: (10000, 10)
```

```
In [31]: Y_prediction
```

```
Out[31]: array([[2.23211202e-04, 5.51889316e-05, 1.63338453e-04, ...,
            8.09079211e-05, 7.35946838e-03, 4.46087643e-06],
            [7.68834725e-03, 3.08940606e-03, 9.68137073e-08, ...,
            2.80909296e-10, 9.89222050e-01, 6.19801526e-08],
            [2.80735425e-05, 0.005551280e-01, 1.98803213e-03],
```

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js

```
...,
[8.19458066e-08, 2.57242959e-07, 2.24625575e-04, ...,
 6.04846049e-03, 1.22072993e-06, 1.16995146e-07],
[3.10936980e-02, 5.96647933e-02, 4.47810907e-03, ...,
 5.71781397e-03, 2.54460581e-04, 6.01701288e-07],
[7.11552605e-12, 5.51200578e-13, 9.32100264e-08, ...,
 9.99882936e-01, 1.24371059e-14, 6.78344672e-11]], dtype=float32)
```


```
In [32]: classes_x = np.argmax(Y_prediction, axis=1)
```

```
In [34]: classes_x.shape
```

```
Out[34]: (10000,)
```


```
In [65]: warnings.simplefilter(action='ignore', category=FutureWarning)
X = x_test
fig, axes = plt.subplots(10,10, figsize=(20,20))
fig.tight_layout(pad=0.1)
print("Model Predict/Reality")
for i,ax in enumerate(axes.flat):
    random_index = np.random.randint(10000)
    ax.imshow(X[random_index], cmap='gray')
    ax.set_title(labels[int(classes_x[random_index])]+ "/" + labels[int(y_test[random_index])])
    ax.set_axis_off()
```

Model Predict/Reality



In [52]:

```
plt.figure(figsize=(20, 20))
sns.heatmap(confusion_matrix(y_test, classes_x), annot=True)
plt.show()
```



ARIA HASSANALI ARAGH . 99222032